

# Aniket Anand Deshmukh

---

## CONTACT INFORMATION

*Google Scholar:* [shorturl.at/cmxJS](https://scholar.google.com/shorturl.at/cmxJS)  
*Web:* <https://aniketde.github.io/>

*Email:* [aniketde@umich.edu](mailto:aniketde@umich.edu)  
*Phone:* (734) 277-3529

## EDUCATION

**University of Michigan (UMich)**, Ann Arbor, MI USA CGPA: 4.0/4.0  
*Ph.D., Electrical and Computer Engineering* *Aug'13 - Jan'19*

- PhD Thesis Projects: Domain Generalization [6], Multitask Learning for Contextual Bandits [11], Simple Regret Minimization for Contextual Bandits [10]
- Graduate Student Instructor: EECS 545 Machine Learning (Fall 2016, Fall 2018)

  
**Indian Institute of Technology Hyderabad (IIT-H)**, India CGPA: 8.63/10.0  
*Bachelor of Technology with Honors, Electrical Engineering* *Jul'09 - Jul'13*

## PROFESSIONAL EXPERIENCE

**Applied Scientist, AWS AI Labs, Amazon**, Santa Clara, CA, USA *Jan'23 - Present*

- Tech Lead, Personalizing Retrieval-Augmented Generation (RAG):
  - Leading the development of a personalized RAG model for an AWS service (Amazon Q) with 10+ team members, leveraging Large Language Models (LLMs) to enhance efficiency and user experience.
  - The proposed approach achieved precise identification of personalization needs in approximately 10% of all queries, successfully boosting performance from 45% to 65% in terms of acceptability metric within these targeted scenarios in Amazon's internal use cases.
- Tech Lead, Next Best Action (NBA):
  - Led the Next Best Action (NBA) project, advancing personalization by recommending actions to boost long-term user engagement and loyalty. Developed a propensity model predicting user actions, with NBA adopted by 10+ AWS customers shortly after launch.
  - Conducted interviews with 20+ external AWS customers and 5+ Amazon teams to understand NBA needs, shaping the project's direction. Collaborated closely with product managers and a team of 10+ engineers to design, model, and launch NBA as an AWS service at Re-Invent 2023.
- Chat for Data Prep using Large Language Models (AWS Low code No Code):
  - Contributed to the "Chat for Data Prep" feature in AWS Sagemaker Canvas, utilizing LLMs for intuitive chat-based data queries, manipulation, and visualization..
  - Designed the automation for query recommendations using LLMs and created an automated testing engine to assess the "Chat for Data Prep" pipeline.
- Mentor and Hiring Lead, Amazon Personalize 2024 Internships and Mentor for 2023 Interns:
  - Spearheaded the intern selection process, from outreach to finalizing hires. Successfully short-listed over 15 candidates from hundreds of applications, conducted interviews, and coordinated the allocation of selected interns to various project teams.
  - Submitted two papers related to LLM and human feedback at ICML with 2023 interns [1, 2]

  
**Senior Applied Scientist, Microsoft Ads**, Mountain View, CA, USA *Mar'19 - present*

- Multi Media Ads (Retrieve images from general purpose corpus for Ad text):
  - Enhanced Multi Media Ads by developing an image-text multimodal model, achieving a 1.5% increase in click-through rate via A/B testing and outperforming OPEN AI's CLIP model by 9.4% in internal evaluations.
  - Coordinated across functions, leading meetings, documentation, and GPU resource requests, and collaborated with managers and engineers to establish an evaluation pipeline and A/B testing framework.
- Image attribute classification for product shopping:
  - Developed and deployed a curriculum learning-based image classification model for attribute classification on Microsoft Ads shopping platform, generating \$2M daily revenue; work published in ECCV 2020 [8].

- Labelling images for the project was expensive and time consuming. Mentored an intern on a state-of-the-art clustering method to streamline image labeling, enhancing team efficiency; findings published in IJCNN 2021 [5].
- Smart Campaigns (AI powered feature to help small businesses manage their advertising):
  - Developed a representation learning model to assess keyword-query similarity based on click data, now aiding over 10,000 advertisers on the Microsoft Ads platform in identifying relevant keywords for their ads.

**Graduate Student Research Assistant, UMich, Ann Arbor, MI, USA**

*Aug'13 - Jan'19*

- Contextual Bandits:
  - Modeled the problem of providing personalized recommendations in the multi-task learning for contextual bandits setting and proved the regret bound [11].
  - Modelled a resource constrained, adaptive sensor selection in interplanetary spacecraft using a novel simple regret minimization framework in contextual bandits. Demonstrated improvement in simple regret of more than 25% over algorithms designed to minimize the cumulative regret [10]. This project got the multiple awards especially the best paper award at the ICML 2019 workshop of "Exploration in RL".
- Domain Generalization (Out of Distribution Learning):
  - Reduced the time complexity of kernel-based domain generalization algorithm from  $O(n^2)$  to  $O(n)$  using kernel approximation technique and proved the upper bound on the approximation error [6].

**Applied Scientist Intern, Microsoft AI & Research, Sunnyvale, CA, USA**

*May'18 - Aug'18*

- Proposed and implemented Word2Vec variant of CDSSM algorithm to improve text embedding of a search query from click feedback.
- Developed self supervised learning for contextual bandits algorithms. Achieved the cumulative reward gains from 0.5% to 46.11 % over Neural-UCB on standard computer vision datasets [12].

**Research Intern, Mitsubishi Electric Research Labs, Cambridge, MA, USA**

*May'16 - Aug'16*

- Proposed a distribution free, graph based approach for semisupervised transfer learning.
- Demonstrated that the additional information added in the form of unlabeled data improves prediction by 5.9% on the Parkinson's Telemonitoring dataset as compared to transfer learning method.

**SKILLS**

*Languages*

Python, SQL, Latex, HTML

*Tools*

PyTorch, Keras, TensorFlow, Scikit-learn, Pandas, Matlab, Azure ML, AWS, etc

**ACADEMIC  
SERVICE**

- Outstanding Reviewer award at NeurIPS 2021, AISTATS 2022 and NeurIPS 2023 *Jun'21 - current*
- Reviewer, NeurIPS, ICML, ICLR, CVPR, AISTATS, JMLR, UAI, AAAI, etc *Jan'18 - current*
- Co-chair, TheWebConf 2023 Decision Making for IR and Recommender Systems *Oct'22 - Apr'23*
- Co-chair, ICLR 2023 Domain Generalization Workshop *Oct'22 - Apr'23*

**SELECTED  
PUBLICATIONS**

1. Subhojyoti Mukherjee, Ge Liu, Aniket Deshmukh, Anusha Lalitha, Yifei Ma, and Branislav Kveton. Experimental design for active transductive inference in large language models. *ICML submission*, 2024
2. Subhojyoti Mukherjee, Anusha Lalitha, Kousha Kalantari, Aniket Deshmukh, Ge Liu, Yifei Ma, and Branislav Kveton. Optimal design for k-way human feedback. *ICML submission*, 2024
3. Aniket Anand Deshmukh, Jayanth Reddy Regatti, Eren Manavoglu, and Urun Dogan. Representation learning for clustering via building consensus. *arXiv preprint arXiv:2105.01289*. Accepted at *Springer Machine Learning Journal*, 2022
4. Jayanth Reddy Regatti, Aniket Anand Deshmukh, Eren Manavoglu, and Urun Dogan. Consensus clustering with unsupervised representation learning. In *2021 International Joint Conference on Neural Networks (IJCNN)*, pages 1–9. IEEE, 2021

5. Gilles Blanchard, Aniket Anand Deshmukh, Urun Dogan, Gyemin Lee, and Clayton Scott. Domain generalization by marginal transfer learning. *Journal of Machine Learning Research*, 22(2):1–55, 2021
6. Jayanth Reddy Regatti, Aniket Anand Deshmukh, Frank Cheng, Young Hun Jung, Abhishek Gupta, and Urun Dogan. Offline rl with resource constrained online deployment. *arXiv preprint arXiv:2110.03165*, 2021
7. Urun Dogan, Aniket Anand Deshmukh, Marcin Machura, and Christian Igel. Label-similarity curriculum learning. In *European Conference on Computer Vision*, pages 174–190. Springer, 2020
8. Aniket Anand Deshmukh, Srinagesh Sharma, James W Cutler, Mark Moldwin, and Clayton Scott. Simple regret minimization for contextual bandits. In *Exploration in RL, ICML 2019 workshop arXiv preprint arXiv:1810.07371*, 2018
9. Aniket Anand Deshmukh, Urun Dogan, and Clay Scott. Multi-task learning for contextual bandits. In *Advances in Neural Information Processing Systems*, pages 4848–4856, 2017